

JSON stores

Ioana Manolescu Silviu Maniu

INRIA Saclay

U. Paris-Sud

M2 Data and Knowledge

Université Paris Saclay

Motivation


- JSON (JavaScript Object Notation) allows to describe nested, potentially heterogeneous data
 - Very flexible
 - Thus, a good idea for NoSQL!
- Much less verbose than XML

Sample JSON document: Twitter

```
{ "results": [  
  { "text": "@twitterapi http://tinyurl.com/ctrefg",  
    "to_user_id": 396524,  
    "to_user": "TwitterAPI",  
    "from_user": "jkoum",  
    "metadata": { "result_type": "popular", "recent_retweets": 109 },  
    "id": 1478555574,  
    "from_user_id": 1833773,  
    "iso_language_code": "nl",  
    "source": "twitter< /a>",  
    "profile_image_url": http://s3.amazonaws.com/twitter/a155\_b\_normal.jpg,  
    "created_at": "Wed, 08 Apr 2009 19:22:10 +0000"},  
    ... truncated ... ],  
  "refresh_url": "?since_id=1480307926&q=%40twitterapi",  
  "results_per_page": 15,  
  "next_page": "?page=2&max_id=1480307926&q=%40twitterapi",  
  "completed_in": 0.031704,  
  "page": 1,  
  "query": "%40twitterapi"}  
}
```

JSON document structure

- Object: collection of (name, value) pairs




- Array: collection of values



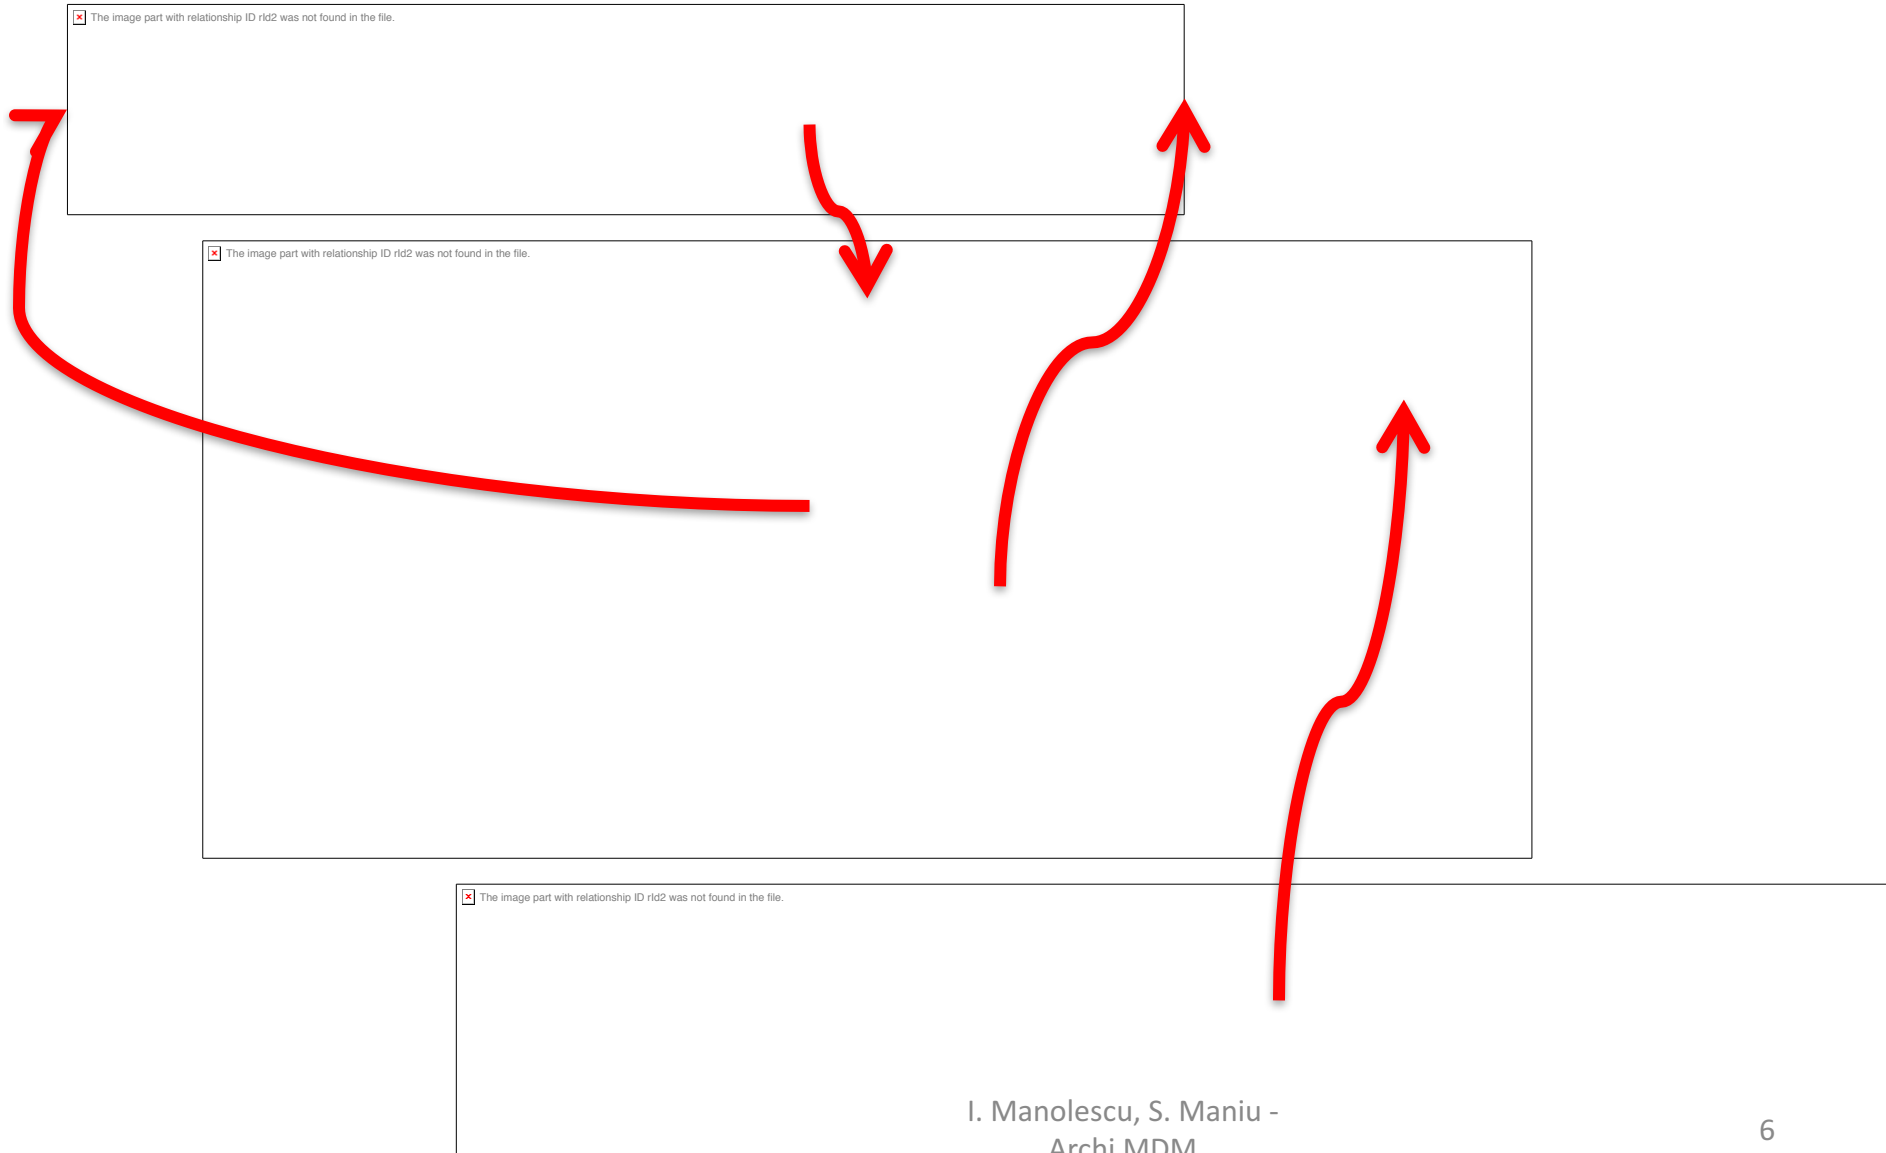
- Arrays and object structure are heterogeneous (no schema)

JSON document structure


- Values (allow nesting):



JSON document structure



MongoDB: a JSON document store

 The image part with relationship ID rld2 was not found in the file.

Computed based on: popularity in search engine results, queries, job offers,
social networks, questions on StackOverflow...

MongoDB architecture

- Fundamentally **client-server**
- The server data can be
 - **Replicated**
 - Several identical copies of the same data
 - To be seen
 - **Partitioned** (sharded)
 - Distributed across the machines of a cluster in order to take advantage of the storage and processing capacity
 - To be seen
- Document processing
 - Selective access
 - Map-Reduce mode
- Uses a special form of JSON, called **BSON** (binary)

MongoDB commands

- Allows **CRUD** commands
 1. **Create** db.col.insertOne(), db.col.insertMany()
 2. **Read** db.col.find(<condition>)
 3. **Update** db.col.updateOne(),
db.col.updateMany(), db.col.replaceOne()
 4. **Delete** db.col.deleteOne(), db.col.deleteMany()

MongoDB storage organization

- **Documents** are stored in **collections** (which may have **indexes**)
- Collections are part of **databases**

```
> mongo myExample           // Creates the database myExample
> db.towns.insert({         // Creates the collection towns
  name: "New York",         // and inserts a document into it
  population: 22200000,
  last_census: ISODate("2009-07-31"),
  famous_for: [ "statue of liberty", "food" ],
  mayor : {
    name : "Michael Bloomberg",
    party : "I"
  }
}
```

MongoDB object IDs


```
> db.towns.find()  
{  
  "_id" : ObjectId("4d0ad975bb30773266f39fe3"),  
  "name" : "New York",  
  "population": 22200000,  
  "last_census": "Fri Jul 31 2009 00:00:00 GMT-0700  
(PDT)",  
  "famous_for" : [ "statue of liberty", "food" ],  
  "mayor" : { "name" : "Michael Bloomberg", "party" : "I" }  
}
```

_id is implicitly added by the system

Each object ID is different

MongoDB object IDs

12 bytes:



Timestamp; client machine ID; process ID;
incremented counter



IDs are unique across machines and databases

MongoDB: information about data

It is possible to ask questions about objects, functions etc.

```
> typeof db
```

```
object
```

```
> typeof db.towns
```

```
object
```

```
> typeof db.towns.insert
```

```
function
```

MongoDB: getting information about data

Calling a function with no parentheses shows the function code

db.towns.insert

```
function (obj, _allow_dot) {  
    if (!obj) { throw "no object passed to insert!"; }  
    if (!_allow_dot) { this._validateForStorage(obj); }  
    if (typeof obj._id == "undefined") {  
        var tmp = obj;  
        obj = {_id:new ObjectId};  
        for (var key in tmp) { obj[key] = tmp[key];}  
    }  
    this._mongo.insert(this._fullName, obj);  
    this._lastID = obj._id;  
}
```

Working with (JavaScript) functions

Typing this into the client shell registers the function:

```
function insertCity(name, population, last_census,  
                    famous_for, mayor_info) {  
    db.towns.insert( {name:name, population:population,  
                    last_census: ISODate(last_census),  
                    famous_for:famous_for,  
                    mayor : mayor_info}  
    );  
}
```

Object implicitly created

Working with functions

Calling the function previously defined:

```
insertCity("Punxsutawney", 6200, '2008-31-01',  
["phil the groundhog"], { name : "Jim Wehrle" })
```

Array

Object

```
insertCity("Portland", 582000, '2007-20-09',  
["beer", "food"],  
{ name : "Sam Adams", party : "D" } )
```

db.towns.find() returns three objects

Searching a MongoDB collection

```
db.towns.find({ "_id" :  
ObjectId("4d0ada1fbb30773266f39fe4") })
```

returns the full object

```
db.towns.find({"_id" :  
ObjectId("4d0ada1fbb30773266f39fe4") }, ) }, { name : 1 })  
{  
  "_id" : ObjectId("4d0ada1fbb30773266f39fe4"),  
  "name" : "Punxsutawney"  
}
```

Searching a MongoDB collection

To exclude an attribute from a result, **set it to 0** in the find parameter

```
db.towns.find({ _id : ObjectId("4d0ada1fbb30773266f39fe4") }, {  
  name : 0 })  
{  
  "_id" : ObjectId("4d0ada1fbb30773266f39fe4"),  
  "population" : 6200,  
  "last_census" : "Thu Jan 31 2008 00:00:00 GMT-0800 (PST)",  
  "famous_for" : [ "phil the groundhog" ]  
}
```

Searching a MongoDB collection

```
db.towns.find(  
    { name : /^P/, population : { $lt : 10000 } },    //selection  
    { name : 1, population : 1 })                    // projection
```

→

```
{ "name" : "Punxsutawney", "population" : 6200 }
```

```
db.towns.find(  
    { last_census : { $lte : ISODate('2008-31-01') } },  
    { _id : 0, name: 1 })
```

→

```
{ "name" : "Punxsutawney" }  
{ "name" : "Portland" }
```

Searching in nested structures

```
db.towns.find(  
  { famous_for : "food" },  
  { _id : 0, name : 1, famous_for : 1 })
```

→

```
{ "name" : "New York", "famous_for" : [ "statue of liberty", "food" ] }  
{ "name" : "Portland", "famous_for" : [ "beer", "food" ] }
```

```
db.towns.find(  
  { famous_for : { $all : ["food", "beer"] } },  
  { _id : 0, name:1, famous_for:1 })
```

→

```
{ "name" : "Portland", "famous_for" : [ "beer", "food" ] }
```

\$all matches an array
containing all the specified
values

Searching in nested structures

Nodes which *must not* have a match of the search condition:

```
db.towns.find(  
  { famous_for : { $nin : ["food", "beer"] } },  
  { _id : 0, name : 1, famous_for : 1 } )  
→  
  { "name" : "Punxsutawney", "famous_for" : [ "phil the groundhog" ] }
```

Paths in conditions:

```
db.towns.find( { "mayor.party" : "I" }, { _id : 0, name : 1, mayor : 1 } )  
→  
{ "name" : "New York",  
  "mayor" : { "name" : "Michael Bloomberg", "party" : "I" }  
}
```

Searching in nested structures

Countries that export **bacon and tasty food**:

```
db.countries.find({ "exports.foods.name" : "bacon", "exports.foods.tasty" :  
true }, { _id : 0, name : 1 } )
```

Countries that export **tasty bacon**:

```
db.countries.find(  
{"exports.foods" : { $elemMatch : { name : "bacon",  
tasty : true }  
}}, { _id : 0, name : 1 } )
```

Matched by:

```
{_id : "us", name : "United States",  
exports : { foods : [ { name : "bacon", tasty : true }, { name : "burgers" } ] }  
}
```

More search operators

\$regex: matches PCRE-compliant regexes within / /

\$ne, \$lt, \$lte, \$gt, \$gte: arithmetics

\$exists, \$all, \$in, \$nin, \$or, \$nor, \$not: logical operators

\$elemMatch

\$size: matches array of given size

\$type: matches if field is of a given type

db.countries.find(

```
{ $or: [ { _id: "mx" }, { name: "United States" } ] },  
{ _id: 1 } )
```

Updates

```
db.towns.update(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") },  
  { $set : { "state" : "OR" } } );  
                                     // updates state
```

```
db.towns.update(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") },  
  { { "state" : "OR" } } );  
                                     // replaces the whole document!
```

```
db.towns.delete(  
  { _id: ObjectID("4d0ada87bb30773266f39fe5") } )  
                                     // deletes the document
```


More operators used in updates

\$set, \$unset (removes the field)

\$inc (increments)

\$pop, \$push, \$pushall for arrays

\$addToSet like push but avoids duplicates

\$pull removes a matching value from an array

\$pullAll removes all matching values

Combining information from several documents

MongoDB does not provide joins!

- In most cases, data that should be used together is stored in the same document(s)

To combine data from several documents, two options:

1. Store the ID of a document within another:

```
original_id = ObjectId()  
db.places.insert({ "_id": original_id,  
                  "name": "Broadway Center",  
                  "url": "bc.example.net" }) db.people.insert({  
"name": "Erin",  
                  "places_id": original_id,  
                  "url": "bc.example.net/Erin" })
```

Collection and database
are not specified!
Either consider the same
or extra logic is needed.

Combining information from several documents

To combine data from several documents, two options:

2. Use a **DBRef** instance, which is recognized as a pointer to another document:

```
{ "$ref" : <value>, "$id" : <value>, "$db" : <value> }
```

"\$ref" points to the collection; "\$id" points to the document; "\$db" points to the database

```
{ "_id" : ObjectId("5126bbf64aed4daf9e2ab771"),    // ..  
  "creator" : { "$ref" : "creators", "$id" :  
                ObjectId("5126bc054aed4daf9e2ab772"), "$db" : "users" } }
```

Recognized in some language drivers, not in all

E.g. Java, Python, Perl, PHP OK

E.g. C, C++, Scala: not supported (yet)

Global processing with custom code

One can always define a JS function and run it

```
db.towns.find(  
  ( function() { return this.population > 60000;} );  
  // runs the function over all the towns
```

Fails if one town has no population!

Contrast with XML/XQuery

"OK for extra, not OK for missing"

Indexing MongoDB data

Indexes can be built on a collection calling `collection.createIndex({attr...})`

1. **B-trees** for exact and inequality search

- May be built on a single attribute (simple) or several attributes (compound)
- `Collection.createIndex({"name": 1})` ; *// for ascending order, otherwise use -1*
- B-tree index automatically built on `_id`

2. **Multikey indexes** allow indexing on an array attribute

- Built by default when one requires the indexing of an array attribute

Indexing MongoDB data

- 3 **Geospatial indexes: 2d** (planar geometry based on x,y), **2d sphere** (latitude, longitude)
 - Operator: **\$near, \$nearsphere** (coordinates) returns the top k closest documents to the given coordinates
 - Operators: **\$geoWithin, \$geoIntersects** (JSON rectangle)
- 4 **Text indexing** for full-text search
 - `db.reviews.createIndex({ comments: "text" })`
 - `db.collection.createIndex({ "$**": "text" })`

Inspect indexes: `db.system.indexes.find()`

Searching with and w/o an index

Without an index:

```
db.phones.find({display: "+1 800-5650001"}).explain()
```

```
{ "cursor" : "BasicCursor",  
  "nscanned" : 109999,  
  "nscannedObjects" : 109999,  
  "n" : 1,  
  "millis" : 52,  
  "indexBounds" : { }  
}
```

With an index on display:

```
db.phones.find({ display: "+1 800-5650001" }).explain()
```

```
{ "cursor" : "BtreeCursor display_1",  
  "nscanned" : 1,  
  "nscannedObjects" : 1,  
  "n" : 1,  
  "millis" : 0,  
  "indexBounds" : {  
    "display" : [ [ "+1 800-5650001", "+1 800-5650001" ] ] }  
}
```

MapReduce processing

ound in the file.

- Execution order:
1. Query
 2. Map
 3. Reduce
 4. [finalize] to wrap up reducer results, e.g. to take the max among all
 5. [output]

Replica sets

Duplication (replication) to prevent against server failure and data loss

Example (three servers): mkdir ./mongo1 ./mongo2 ./mongo3

Create replication set:

```
mongod --replSet book --dbpath ./mongo1 --port 27011 --rest
```

```
mongod --replSet book --dbpath ./mongo2 --port 27012 --rest
```

```
mongod --replSet book --dbpath ./mongo3 --port 27013 --rest
```

Then in one of the servers initialize replication set:

```
mongo localhost:27011
```


```
> rs.initiate({_id: 'book',  
              members: [ {_id: 1, host: 'localhost:27011'},  
                          {_id: 2, host: 'localhost:27012'},  
                          {_id: 3, host: 'localhost:27013'} ] })
```

Then one server will output **[rs Manager] replSet PRIMARY**
while two will output **[rs sync] replSet SECONDARY**

Replica sets

The servers held a vote to determine who is the master (primary); the two others are replicas ("secondary")

By default, applications read/write through the primary, who pushes updates to secondary servers




If the master is considered dead, there are new elections

- *Only succeed if more than half of the original replication set votes*
- Operations attempted on a demoted (dead) master are lost
- A write is considered successful only if $>$ half of the replicas "saw" it

Replica sets

By default, write to the primary, who pushes updates to secondary servers



MongoDB recommends an *odd number of server in a replica set*, to allow a majority in case of network failure. **One arbiter** may be added to the replica set

- Strong consistency on read
- Resistance to some partitioning

MongoDB sharding

Sharding = partitioning

1 shard = 1 fragment

- To distribute a very large collection across several servers

Sharding is logically *on top of replication*

- Each shard server may participate to a replica set



MongoDB sharding


Roles:

- **Shard** (shard server): stores a collection fragment
- **Config server(s)**: store(s) information on which shard has what
- **Single point of entry**: mongos

The image part with relationship ID rid2 was not found in the file.

MongoDB sharding

- A data collection is partitioned into **chunks** based on the value of a **shard key**
- Each chunk covers a key range



- 1 shard = a set of chunks
- Mongos routes writes to the appropriate chunk based on the shard key value
- Chunks are split when they grow beyond a fixed chunk size (64 MB default, can be split)
- MongoDB migrates chunks across shards for load balancing

MongoDB sharding

// starting the shard servers:

```
mkdir ./mongo4 ./mongo5
```

```
mongod --shardsvr --dbpath ./mongo4 --port 27014
```

```
mongod --shardsvr --dbpath ./mongo5 --port 27015
```

// starting the config server:

```
mkdir ./mongoconfig
```

```
mongod --configsvr --dbpath ./mongoconfig --port 27016
```

// starting mongos connected to the config

```
mongos --configdb localhost:27016 --chunkSize 1 --port 27020
```

// talking to mongos to configure sharding:

```
mongo localhost:27020/admin
```

```
> db.runCommand( { addshard : "localhost:27014" } ) -> { "shardAdded" : "shard0000", "ok" : 1 }
```

```
> db.runCommand( { addshard : "localhost:27015" } ) -> { "shardAdded" : "shard0001", "ok" : 1 }
```

```
> db.runCommand({ enablesharding: "test"}) -> { "ok" : 1 }
```

```
> db.runCommand( { shardcollection : "test.cities", key : {name : 1} } ) // { "collectionsharded" :  
"test.cities", "ok" : 1 }
```

Another JSON store: CouchDB

<http://db-engines.com/en/system/CouchDB%3BMongoDB>

Feature	CouchDB	MongoDB
Since	2005	2009
Ranking	#23 overall, #2 document store	#4 overall, #1 document store
From	Apache Software	MongoDB Inc
APIs and other access methods	RESTful HTTP/JSON API	proprietary protocol using JSON
Replication methods	Master-master replication Master-slave replication	Master-slave replication
MapReduce	yes	yes
Consistency concepts	Eventual Consistency	Eventual Consistency Immediate Consistency
Foreign keys	no	no

More document stores: http://db-engines.com/en/ranking_trend/document+store

Concurrency control in CouchDB

- Update granularity = document
- To change a document's attribute, rewrite the document!
- Multi-Version Concurrency Control: some requests may return "old" versions but they each return a version that was valid at some point

 The image part with relationship ID rid2 was not found in the file.

Synchronization in CouchDB

Each server can work independently

Incremental replication can be set to run in the background

Synchronization is on demand between any pair of servers

Diverging changes are flagged as **conflicts**; conflict resolution policy must be specified. One document version wins, the other is considered older.

